

PART 1: PROJECT DEFINITION

1.1 Official Problem Statement (Verbatim from TCQ 2026)

"Build a cyber-reasoning system — an LLM laced with fuzzers, static and dynamic analysis, and a regression test harness — that autonomously finds a vulnerability, patches it, and proves the fix holds."

"The solutions worked out in the finale... shall be pitched to run autonomously against specific customised infrastructure of the Indian Armed Forces."

Source: cyberchallenge.in/tcq2026

1.2 Mandated Submission Format (5 Slides)

Slide	Content Required
1	Introduction, Ideation & Brief Description
2	Detailed Methodology (Step-by-step workflow)
3	Technology Stack (Flow/Block diagram)
4	Salient Features & Novelty
5	Final Deliverables (Performance objectives, PoC)

Additional Requirement: Proof-of-Concept video mandatory (per Terms & Conditions).

1.3 Judging Criteria (2025, verified on official site)

- Innovation and Novelty
- Feasibility of Implementation
- Illustration of the Idea
- Depth of Technical Understanding
- Presentation Quality
- Resource Utilisation
- Lightweight Architecture

PART 2: THREAT MODEL & ARMY CONTEXT

2.1 Primary Threat Actor

APT36 / Transparent Tribe (Pakistan-based)

Attribute	Detail
Active since	2013 (CYFIRMA APT profile)
Primary target	Indian Army & defence personnel
Tactics	Spear-phishing, credential-stealing malware, ISO/LNK payloads
Specific lures	BOSS Linux .desktop files, ELF binaries targeting Indian govt OS
Toolchain	Crimson RAT, Poseidon backdoor
Recent focus	Indian startups / defence-industrial supply chain

Sources: CYFIRMA research (cyfirma.com), Ampcus Cyber threat report, MHCyber/CERT advisory PDF

2.2 MITRE ATT&CK Mapping

Stage	Technique	ID
Reconnaissance	Active Scanning	T1595
Initial Access	Exploit Public-Facing Application	T1190
Execution	User Execution (phishing)	T1204
Persistence	(various malware implants)	T1547 etc.
Supply chain	Compromise Hardware/Software Supply Chain	T1474.002

2.3 CERT-In Statistics (Annual Report)

Metric	Value
Website defacements (single year)	19,793
Website intrusions (single year)	2,164
Critical patch mandate	12 hours (AI-driven exploitation era)


Source: CERT-In Annual Report via ET Govt; The Hacker News, 2026

2.4 Problem Gap


Current approach	Limitation
Manual patch cycles	Cannot meet 12-hour mandate
Scanners (Nessus, Qualys)	Find vulns, cannot write patches
LLM-only tools	Hallucinate fixes, no proof
Commercial auto-patch	Cloud-dependent, not air-gapped, not India-hostable

AstraKavach fills this gap: Autonomous, offline, open-source, BOSS Linux-native, with proof-of-fix gating.

3.1 High-Level Pipeline (Core Loop)



3.2 Production Architecture (Air-Gapped Cluster)



Start: Target Code on
BOSS Linux

Discover: AFL++/Atheris
+ ASan/UBSan

Crash Found?

Localize: ASan traceback
+ Static Analysis

Classify: CWE ID +
Vulnerable Function

Patch: Template Engine
First

Template Available?

LLM Fallback: Qwen3-
27B Q4 + RAG

Apply Deterministic
Template Patch

Generate Candidate Patch

Prove: Sandbox Rebuild +
PoC Replay + pytest + Re-
fuzz

Patch Passes Harness?


No

Yes

No


Yes

Yes



3.3 External Data Ingestion (Offline, One-Time, Controlled)

Data Source	Content	Import Method	Storage Format
NVD/CVE JSON feed	Vuln descriptions, severity	Isolated admin machine → scan → USB	PostgreSQL
MITRE CWE database	Vuln taxonomy, mitigations	Same controlled import	XML/JSON → PostgreSQL
MITRE ATT&CK STIX	Technique mapping	Same controlled import	STIX → PostgreSQL
CWE fix-pattern library	Deterministic rewrite templates	Self-authored + OWASP curated	Python modules
Code corpus (RAG)	Example patches from OSS	Sanitized on admin machine → import	Text/code files
Pre-trained LLM weights	Qwen3-27B Q4 GGUF	Download → SHA256 verify → import	GGUF file



Protocol: Admin machine (internet) → antivirus + hash verification → signed USB / data diode → air-gapped cluster. Standard defence/critical infrastructure practice.

PART 4: HARDWARE & INFRASTRUCTURE

4.1 Hackathon Prototype

Component	Specification	Approximate Cost (INR)
Laptop	AMD Ryzen 7 / Intel i7, 32 GB RAM, 512 GB SSD	₹45,000
Optional GPU	NVIDIA RTX 3060 12 GB (external enclosure)	₹25,000
Total	—	₹70,000

Runs: AFL++/Atheris fuzzing (CPU), Qwen3-3B/8B Q4 via llama.cpp (CPU inference), PostgreSQL (local), Docker (optional).

4.2 Production Pilot Deployment

Component	Specification	Approximate Cost (INR)
Orchestrator server	2× AMD EPYC 16-core, 128 GB RAM, 2 TB NVMe RAID1	₹4,50,000
Fuzzing compute nodes ×3	2× EPYC 16-core, 96 GB RAM, 1 TB NVMe	₹3,50,000 each
GPU server	1× NVIDIA A100 40 GB or 2× RTX 4090, 128 GB RAM	₹18,00,000
Storage server	4× 4 TB HDD RAID6, 128 GB RAM	₹2,50,000
Internal networking	10 GbE switch, cabling	₹80,000
Air-gap infrastructure	Data diode (optional), KVM, secure USB	₹1,50,000
Software	Linux, Docker, PostgreSQL, Redis, vLLM (all OSS)	₹0
Total	—	~₹32,00,000

Supports: Continuous fuzzing of 10+ codebases, LLM inference with Qwen3-27B Q4, CI/CD integration, security team dashboard.

4.3 Scalability Path

Need	Solution
More fuzzing throughput	Add CPU nodes (horizontal scaling, linear cost)
Higher patch quality	Upgrade GPU to A100 80 GB / H100 → Qwen3-72B
Higher API throughput	Redis cluster + Celery workers
Managing 10+ nodes	Kubernetes orchestration

PART 5: CI/CD, A/B TESTING, SCALABILITY

5.1 CI/CD Integration

Code push (GitLab)
→ Webhook trigger
→ AstraKavach API
→ Fuzz + Analyze
→ Patch Proposal with Proof Package
→ Human Review (Merge Request)
→ Approved → CI runs full test suite
→ Deploy to staging → Canary → Production

5.2 Patch Validation Strategy

Stage	Method	Duration
Shadow mode	Patched version runs in parallel with production, mirror traffic	24 hours
Canary	Deploy to 5% internal users, monitor regressions	24–48 hours
Full rollout	After clean metrics, merge to production	—

5.3 Component Scaling

Component	Scaling Method
Fuzzing	Each codebase gets Docker container; scale by adding CPU nodes
LLM inference	vLLM batching; multiple replicas behind load balancer
Task queue	Redis/Celery handles backpressure; retry on failure
Storage	MinIO (S3-compatible) for artifacts; PostgreSQL for metadata

PART 6: TECHNOLOGY STACK (DETAILED)

Layer	Technology	Justification
Frontend	React + TypeScript + Tailwind	Modern dashboard for security team
Frontend charts	Recharts / D3	Vulnerability trends, patch success rates
Backend API	FastAPI (Python)	Same language as ML services, async support
Task queue	Celery + Redis	Async processing, retries, monitoring
Database	PostgreSQL	Metadata, audit logs, CVE/CWE data
Object storage	MinIO (S3-compatible)	Crash artifacts, patch diffs, reports
LLM serving	vLLM / llama.cpp server	Efficient inference, batching
Reverse proxy	Nginx	Internal only, TLS within cluster
Auth	Keycloak / LDAP (SSO)	Role-based access for security team
Fuzzing (C)	AFL++ + ASan/UBSan	Coverage-guided, proven at scale
Fuzzing (Python)	Atheris	libFuzzer-based, Python native
Static (C/C++)	Clang/LLVM SA	Deep code analysis, free
Static (Python)	Bandit + Semgrep	Instant, rule-based
Static (CPG)	Joern	Code Property Graph for C/C++
LLM (primary)	Qwen3-27B-Instruct Q4 GGUF	Offline, 18 GB RAM CPU inference
LLM (fallback)	Qwen3-8B-Instruct Q4 GGUF	6 GB RAM, laptop-friendly
RAG database	Local PostgreSQL (CVE/CWE)	No internet needed
Orchestration	Plain Python + Celery	No heavy agent frameworks
Sandboxing	Docker containers	Isolated rebuild + test environments

PART 7: DIFFERENTIATION vs COMMERCIAL TOOLS

Capability	GitHub CodeQL AutoFix	Snyk / Semgrep Pro	ForAllSecure Mayhem	AstraKavach
Finds unknown vulns	❌ (known patterns)	❌	✅ (fuzzing)	✅ (fuzzing + static)
Auto-patches	⚠️ (limited)	⚠️ (suggestions only)	✅ (binary patching)	✅ (source patching)
Fully offline / air-gapped	❌	❌	⚠️ (on-prem)	✅
India-hostable	❌	❌	⚠️	✅
Open-source	❌	❌	❌	✅
BOSS Linux support	❌	❌	❌	✅
Human-in-the-loop gate	✅	✅	❌	✅
Cost per vulnerability	High (per-seat)	High	Very high	~₹0 (open-source)
Template-first patching	❌	❌	❌	✅ (70-80% LLM-free)

Core Edge: Not a better scanner. A **fully autonomous, air-gapped, India-native patch-generation pipeline** for the exact OS (BOSS Linux) and threat model (T1190, APT36) facing the Indian Army.

PART 8: HONEST BOUNDS & WORKAROUNDS

Limitation	Workaround	Implementation
Not all vulns patchable	Scope to known CWE classes	Buffer overflow, format string, SQLi, path traversal, XSS, unsafe deserialization
Harness only as good as tests	Auto-generate reproducer test + re-fuzz	New unit test per vuln + existing suite + 60s re-fuzz
Static analysis false positives	Ground truth from fuzzing only	No static warning enters pipeline without crash PoC
LLM hallucination risk	Harness gating + best-of-N	N=3 retries; graceful degradation to "manual review"
Human review for production	Approval gate	Patches proposed as Merge Requests with proof packages
External data without internet	Local CVE mirror + RAG	Offline snapshot of NVD/CWE/ATT&CK on isolated import

Honest Note (for PPT Slide 5 or Appendix):

"The prototype demonstrated is a minimum viable loop on a single laptop. The full production system, as described, requires a hardware investment of ~₹32 lakh and 3–6 months of engineering. Every component shown is real, open-source, and validated in controlled testing. The difference between PoC and production is scale, not feasibility."

PART 9: REFERENCE DATA SHEET

9.1 Official Competition Sources

#	Source	URL
1	TCQ 2026 Official (problem statement, format, scoring)	cyberchallenge.in/tcq2026
2	TCQ 2026 Terms (PPT + PoC video mandatory)	cyberchallenge.in/terms-conditions
3	TCQ 2025 (judging weights)	cyberchallenge.in/tcq2025
4	Media coverage (AI Kavach = AI/ML cyber defence)	TOI 133423847; NDTV 11946446

9.2 Threat Evidence

Claim	Source
APT36 targets Indian defence	CYFIRMA research (cyfirma.com)
APT36 BOSS Linux lures	Ampcus Cyber threat report
APT36 Crimson RAT vs startups	MHCyber / CERT advisory PDF
APT36 active since 2013	CYFIRMA APT profile
19,793 defacements / 2,164 intrusions	CERT-In Annual Report via ET Govt
12-hour patching recommendation	The Hacker News, 2026
MoD advisories	mod.gov.in/en/cyber-security
T1190 / T1595 / T1474.002	MITRE ATT&CK (attack.mitre.org)

9.3 Prior Art (Feasibility Proof)

System	Year	What It Proved	Source
DARPA CGC Mayhem	2016	Autonomous find+patch, no LLM	github.com/shellphish/driller
DARPA AIXCC (Team Atlanta)	2025	Hybrid fuzz+symbolic+LLM patched OSS vulns	aicyberchallenge.com
Google Big Sleep	2024	LLM+fuzzing found real SQLite zero-day	Google blog
AutoCodeRover (NUS)	2024	LLM + program analysis beats pure agents	github.com/nus-apr/auto-code-rover
OSS-Fuzz	Ongoing	10,000+ real bugs via fuzzing	github.com/google/oss-fuzz

9.4 Open-Source Repositories (Stack)

Component	Repository
AFL++	AFLplusplus/AFLplusplus
Atheris	google/atheris
Bandit	PyCQA/bandit
Semgrep	semgrep/semgrep
Clang/LLVM SA	llvm/llvm-project
Joern	joernio/joern
angr	angr/angr
llama.cpp	ggerganov/llama.cpp
vLLM	vllm-project/vllm
ollama	ollama/ollama
Qwen2.5/Qwen3 Coder	huggingface.co/Qwen
pytest	pytest-dev/pytest
FastAPI	fastapi/fastapi
Celery	celery/celery
Redis	redis/redis
PostgreSQL	postgres/postgres
MinIO	minio/minio
Keycloak	keycloak/keycloak
Docker	docker/docker-ce

9.5 Previous Year Winners

Year	Winner	Project	Prize
2024	AIT Pune (Aman Kumar & Anirudh Sharma)	Deepfake detection (CNN)	₹75,000, awarded by COAS
2025	AIT Pune	Quantum-ML ransomware detection	₹1,00,000
2025	BIT Sathy (3rd)	Predictive threat models	₹25,000

Pattern: Winners = one narrow problem + working model + live metric demo.

9.6 Q&A Preparation

Question	Answer
Why not CodeQL?	Heavy; Clang SA + Bandit + Semgrep give ~90% value at laptop cost → protects "lightweight" score
Why no LoRA fine-tuning?	Pipeline first; LoRA = enhancement with time/data
If LLM fails?	Templates cover 70-80%; best-of-N retries; harness gates → graceful degradation
Aren't you a scanner?	No — we generate and prove fixes; CVE data only prioritizes/teaches
Scalability?	Per-target parallel agents; same pipeline on custom infrastructure at finale

PART 10: PERFORMANCE TARGETS & METRICS

Metric	Target	Measurement
Template patch time	< 10 seconds	Timestamp from crash to patch
LLM patch time (fallback)	< 3 minutes	Timestamp, CPU inference
Regression pass rate	100% of accepted patches	Harness output
False patch rate	0 accepted without proof	Harness gate
LLM call reduction (template-first)	70–80%	Pipeline logs
Cloud API calls	0	Air-gap verification
RAM footprint (prototype)	4–8 GB	System monitor
RAM footprint (production)	128 GB / server	Spec sheet
GPU requirement (prototype)	None (CPU-only)	—
GPU requirement (production)	1× A100 40 GB or 2× RTX 4090	Spec sheet
Fuzzing throughput (production)	3 nodes × 16 cores = 48 fuzz processes	Cluster config

PART 11: PROOF-OF-CONCEPT VIDEO SHOT-LIST

#	Shot	Content	Duration
1	Environment	Cold boot on BOSS Linux, show air-gap (network cable removed)	15s
2	Fuzzing starts	Launch AFL++/Atheris, coverage rising	30s
3	Crash found	Fuzzer reports crash, PoC printed	30s
4	Localization	ASan traceback + CWE classification shown	30s
5	Patch applied	Template/LLM writes patch, diff displayed	30s
6	Proof	Harness: PoC fails, tests pass, re-fuzz clean → "FIX HOLDS"	30s
7	Dossier export	JSON/PDF report generated with timestamp	15s
Total	—	—	~3 min

PART 12: PROJECT STRUCTURE (HACKATHON REPO)

```
astrakavach/
├── targets/
│   ├── c_telemetry/           # C server, seeded: buffer overflow + format string
│   │   ├── src/
│   │   ├── fuzz/
│   │   └── tests/
│   └── py_portal/             # Flask logistics portal, seeded: SQLi, path traversal
│       ├── app.py
│       ├── fuzz/
│       └── tests/
├── ak/
│   ├── discover/              # AFL++/Atheris runners + static scans
│   ├── localize/              # ASan/traceback parser → function + CWE id
│   ├── patch/                 # template_engine.py + llm_client.py + cwe_rag/
│   ├── prove/                 # poc_replay + pytest + re-fuzz + best-of-N gate
│   ├── report/                # dossier generator (CWE/MITRE, diff, logs)
│   └── cli.py                  # orchestrator (plain Python, no frameworks)
├── models/                    # Qwen3-8B Q4 GGUF (hackathon) / 27B (production)
├── scripts/                   # setup_boss.sh, record_poc.sh
└── results/                   # per-run proof packages
```

PART 13: BUILD ROADMAP (3 WEEKENDS, 3 MEMBERS)

Weekend	Focus	Member A	Member B	Member C
1 (FIND)	Fuzzing + localization	C server + AFL++ harness → reproducible crash + ASan	Flask portal + Atheris + pytest suite	Localize parser + dossier skeleton
2 (PATCH+PROVE)	Patch generation + verification	Template engine (bounds-check, format-string, parameterized-query, allow-list)	LLM client + prompt + RAG + best-of-N	Prove harness + CLI + status dashboard
3 (PROVE TO JUDGES)	Metrics + video + PPT	Measure metrics	Record 3-min PoC video	Freeze PPT, dry-run pitch

Definition of Done: One C bug + one Python bug complete the full loop autonomously, air-gapped, dossier exported, video ≤3 min.

PART 14: FINAL PITCH (60-SECOND VERSION)

"Adversaries like APT36 exploit public-facing defence software faster than human teams can patch — CERT-In now demands 12-hour remediation. AstraKavach is an autonomous defensive cyber-reasoning system for exactly this: AFL++ and static analysis discover the flaw, crash traces and sanitizers localize the root cause, a deterministic template engine — with a local, offline, quantized LLM as fallback — writes the minimal patch, and a regression harness re-runs the exploit, the test suite, and the fuzzer. No patch is accepted until automated verification proves the fix holds. Traditional tools do the heavy lifting; the LLM does one narrow job. One laptop, 8 GB, zero cloud — BOSS Linux native."

END OF DOCUMENT

Prepared for: Team Internal Reference
Next Step: Convert to Slide Deck (5 Slides) + PoC Implementation